



ADA RESOURCE LIBRARY

Content Moderation

BEGINNER'S GUIDE

Accurate Data. Trusted Annotation. Real AI Results.

1. Introduction to Content Moderation at ADA

Welcome to the ADA Content Moderation Beginner Guide. This document introduces the principles, workflow, and quality standards that every ADA content moderator follows when reviewing data for AI training and platform safety projects.

Content moderation sits at the intersection of human judgment and machine learning: the labels you apply directly shape how AI systems recognize and respond to harmful, sensitive, or policy-violating content. At ADA, we believe accurate, consistent, and well-documented moderation decisions are what separate trustworthy AI training data from unreliable data.

This guide walks you through what content moderation is, the categories of content you will encounter, the step-by-step review process, and the quality standards ADA expects from every moderator.

2. What Is Content Moderation?

Content moderation is the process of reviewing content — text, images, audio, or video — against a defined set of policy guidelines and applying the correct label, flag, or action.

Unlike straightforward annotation tasks such as part-of-speech tagging, content moderation requires contextual judgment: the same word or image can be acceptable in one context and a violation in another.

In the AI industry, content moderation data is used to train and evaluate models that detect hate speech, misinformation, graphic violence, spam, and other harmful categories before that content ever reaches real users. Every decision a moderator makes becomes a training signal, so precision and consistency matter as much as speed.

3. General Moderator Rules

Every ADA moderator follows these baseline rules on every project, regardless of client or category:

- Apply the official guideline document, not personal opinion, every time a decision is made.
- Stay consistent: identical or near-identical content should always receive the same label.
- Read full context before deciding; a sentence, image, or clip can change meaning entirely depending on the surrounding material.
- When a case feels ambiguous, escalate it. Guessing is never acceptable.
- Document the reasoning behind every flagged or escalated decision in the notes field.

- Protect your own wellbeing. Sensitive material can be emotionally taxing; ADA encourages scheduled breaks and supports moderators who need to step away from difficult content.
- Treat every file as confidential. Content reviewed during a project must never be shared, discussed outside the project, or saved locally.

4. Content Moderation Categories

ADA moderators typically work across the following policy categories. Each client project provides its own detailed guideline document, but the categories below represent the industry-standard starting point:

- Hate Speech & Discrimination — content that attacks, demeans, or incites hostility against a person or group based on protected characteristics such as race, religion, gender, nationality, or disability.
- Violence & Graphic Content — depictions or descriptions of real-world violence, including content that glorifies, threatens, or incites physical harm.
- Sexual & Adult Content — material that is sexually explicit or inappropriate for general audiences. Any content involving minors must be escalated immediately rather than labeled.
- Harassment & Bullying — targeted attacks, threats, or demeaning content directed at a specific individual or group.
- Spam & Misleading Content — deceptive, manipulative, or repetitive content designed to mislead users or game a platform's systems.
- Misinformation & Disinformation — false or misleading claims presented as fact, particularly around health, elections, or public safety.
- Self-Harm & Dangerous Behavior — content that promotes, instructs, or glorifies self-harm, suicide, or other dangerous activities.
- Illegal Activity & Regulated Goods — content facilitating illegal transactions, regulated substances, or other unlawful activity.

Important: any content suggesting imminent risk to a real person's safety (self-harm, threats of violence, child exploitation) must be escalated immediately through the emergency escalation path defined in your project guidelines — not simply labeled and left in queue.

5. Step-by-Step Moderation Process

1. Receive the content item along with any available context (platform, surrounding conversation, metadata).
2. Review the item against the current version of the client's policy guideline document — guidelines are updated frequently, so always confirm you are working from the latest version.

3. Identify the applicable category or categories.
4. Determine the correct action: approve, flag for review, label by category, or escalate.
5. Record a short, clear rationale explaining the decision.
6. Submit the item for the next stage of the workflow, which typically includes a QA review pass.

6. Severity Levels & Escalation

Most ADA projects use a four-tier severity scale:

- Low — minor policy friction with no real-world harm risk (e.g. mild rule-breaking).
- Medium — clear policy violation requiring action but no urgent safety risk.
- High — violation with potential for real harm, requiring prompt review by a senior moderator.
- Critical — immediate risk to a person's safety or wellbeing (self-harm, credible threats, child safety). Critical cases must be escalated immediately through the designated emergency path — never delayed for batch review.

When in doubt about severity, always escalate to the next tier up rather than down.

7. Common Mistakes & How to Avoid Them

- Applying personal beliefs instead of the official guideline — always default to the written policy, even when you personally disagree with it.
- Skipping context — reviewing content in isolation, without checking surrounding conversation or metadata, is the leading cause of mislabeling.
- Over-flagging or under-flagging — both reduce data quality. Calibration sessions exist specifically to correct this drift.
- Inconsistency across similar items — keep a mental or written log of edge cases you've seen before, and apply the same logic each time.
- Leaving the rationale field blank — QA reviewers and clients rely on your documented reasoning to audit decisions and refine guidelines.

8. Quality Expectations & Review Criteria

ADA holds every content moderation project to the following standards:

- Accuracy — moderators are expected to maintain at least 95% agreement with QA reviewers and the gold-standard answer set.

- Consistency — repeated or near-duplicate items must receive matching labels across the full dataset.
- Documentation — every flagged, escalated, or rejected item must include a clear, written rationale.
- Turnaround — moderators are expected to meet the throughput targets set for each project while never sacrificing accuracy for speed.
- Calibration — moderators participate in regular calibration sessions where sample decisions are reviewed against the gold standard to identify and correct any drift.

Moderators who fall below quality thresholds receive additional guideline review and coaching before returning to live project work — ADA's goal is always to build skill, not simply to flag underperformance.