



Prompt Engineering Basics for Annotators

As AI clients increasingly task annotators with evaluating, rating, and refining prompts for large language models, a working understanding of prompt engineering has become a core annotation skill. This section introduces the fundamentals every ADA annotator should know before working on prompt-and-response evaluation tasks.

Accurate Data. Trusted Annotation. Real AI Results.

1. What Is Prompt Engineering?

Prompt engineering is the practice of designing and refining the instructions given to a language model so that it produces accurate, relevant, and well-structured output. For annotators, this does not usually mean writing prompts from scratch — it means understanding what makes a prompt effective, so you can judge whether a model's response actually followed it.

2. Why It Matters in Annotation Work

Many ADA projects involve annotators rating or comparing model outputs against the prompts that generated them. A response can only be judged fairly once the annotator understands what the prompt was actually asking for — its explicit instructions and its implied intent. Weak prompt literacy leads to inconsistent ratings and lowers overall dataset quality.

3. Core Principles of an Effective Prompt

- **Clarity & Specificity** — precise, unambiguous instructions with no room for misreading.
- **Context & Role Setting** — background information, persona, and intended audience.
- **Output Format Instructions** — length, structure, language, and tone of the expected answer.
- **Constraints & Negative Instructions** — explicit limits on what the response should not do.
- **Examples (Few-Shot Guidance)** — sample input/output pairs that demonstrate the desired result.

4. Common Prompt Patterns You Will Encounter

- **Zero-shot prompts** — no examples provided; the model relies on the instruction alone.
- **Few-shot prompts** — one or more examples included to guide the response.
- **Chain-of-thought prompts** — the model is asked to reason step by step before answering.
- **Role-based / persona prompts** — the prompt opens with “You are a...” to set a perspective.
- **System + user prompt structure** — a system message sets the rules; the user message makes the request.

5. How Annotators Evaluate Prompt-Response Pairs

When rating a response, check whether it is:

- **Relevant** — addresses what the prompt actually asked.
- **Accurate** — factually correct and free of fabricated details.
- **Complete** — covers every part of a multi-part instruction.
- **Compliant** — respects the stated format, length, and tone constraints.

- **Safe and unbiased** — free of harmful, offensive, or skewed content.
- **Instruction-following** — follows explicit constraints over implied ones when the two conflict.

6. Common Pitfalls to Flag

- Ambiguous instructions that could reasonably be read more than one way.
- Missing context the model would need in order to answer well.
- Conflicting constraints within the same prompt.
- Leading or biased phrasing that pushes the model toward one answer.
- Prompts that ask for something outside the model's stated scope.

7. Quick Reference Checklist

- Did I identify the explicit instruction and the implied intent?
- Did I check the response against every part of the prompt?
- Did I flag missing context or contradictions in the prompt itself?
- Did I separate “the response is wrong” from “the prompt was unclear”?
- Did I record formatting, tone, and length compliance separately from content accuracy?